

# 1

## Affinity Propagation Clustering – Definition

**Affinity Propagation (AP)** is a clustering algorithm that **finds clusters by sending messages between data points**.

Instead of specifying the number of clusters (like in K-Means), **it automatically decides how many clusters to form** based on the data.

It works by finding “**exemplars**” — representative data points that best describe each cluster.

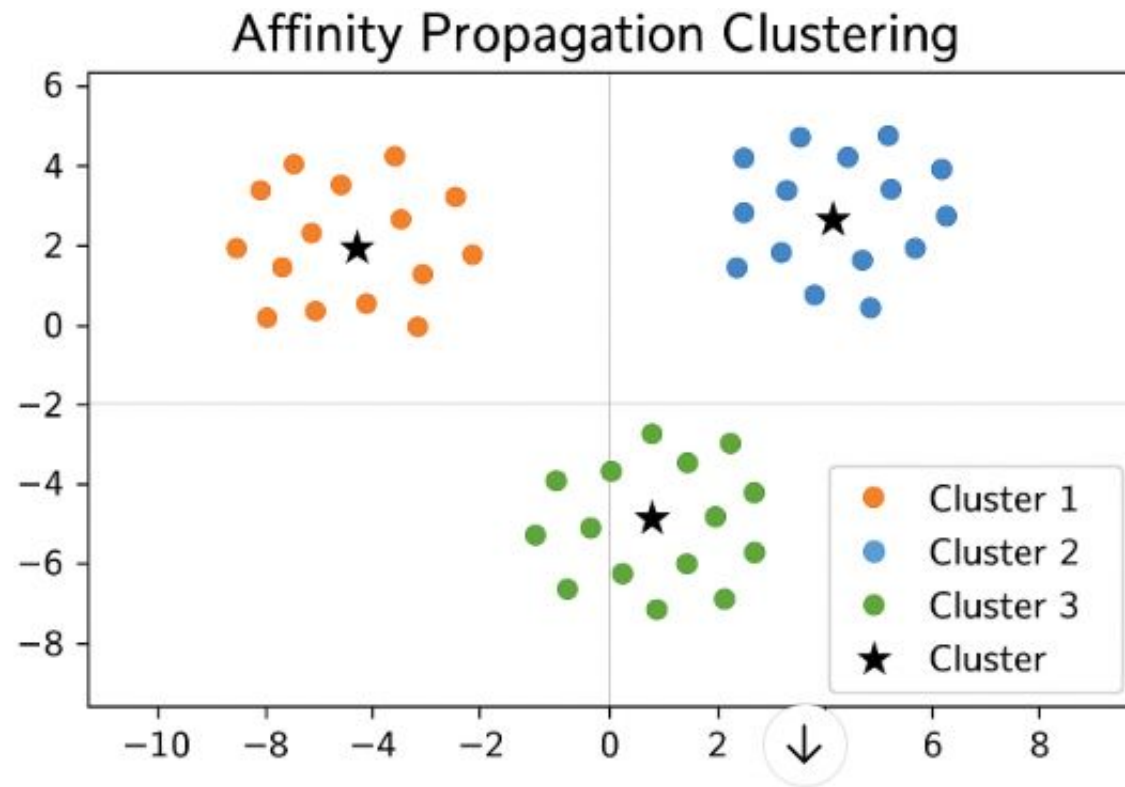
## Advantages

1. No need to specify number of clusters (unlike K-Means).
2. Can find clusters of different sizes and shapes.
3. Works well when there are **clear exemplar (center) points**.
4. Handles **non-metric similarities** (custom similarity measures).


## ✗ Disadvantages

1. **Slow** for large datasets (since it uses all pairwise similarities).
2. Needs careful tuning of parameters ( `damping` , `preference` ).
3. Sometimes results can be **unstable or hard to interpret**.
4. **Memory intensive**, since it stores similarity matrix for all points.

# From sklearn.cluster import AffinityPropagation




## Summary Table

Feature	Description
Type	Unsupervised clustering
Requires #clusters?	 No
Output	Cluster labels and exemplars
Key Parameters	<code>damping</code> , <code>preference</code>
Best for	Medium-sized datasets with clear cluster centers

# 2

## Agglomerative Clustering – Definition

Agglomerative Clustering is a bottom-up hierarchical clustering method.

 It starts by treating each data point as its own cluster,  
then merges the two closest clusters step by step  
until all points belong to one big cluster or until the desired number of clusters is reached.

## Advantages

1. **No need to specify number of clusters initially** (you can decide later by cutting the dendrogram).
2. **Simple and easy to understand.**
3. Works well with **small datasets**.
4. Can handle **non-spherical clusters** (unlike K-Means).



## Disadvantages

1. **Computationally expensive** for large datasets.
2. Once merged, clusters **cannot be split** again.
3. **Choice of distance metric** can affect results.
4. **Sensitive to noise and outliers.**



# from sklearn.cluster import AgglomerativeClustering

## Agglomerative Clustering

Step 1: Each point is a cluster

A ● B C D

Step 2: Merge nearest clusters

AB ● C D E


Step 3: Merge again

ABC ● ● DE

Step 4: Merge all → One big cluster

ABCDE ABCDE

## Agglomerative Clustering







# 3

## BIRCH Clustering – Definition

**BIRCH (Balanced Iterative Reducing and Clustering using Hierarchies)** is a hierarchical clustering algorithm designed to handle **very large datasets efficiently**.

It builds a **tree structure (CF Tree)** that summarizes the data, and then clusters those summaries instead of all individual points — this makes it **fast and memory-efficient**.

## Advantages

1.  **Very fast** — works well on **large datasets**.
2.  **Automatically reduces data size** using CF Tree.
3.  **Memory-efficient**, doesn't need to load all data at once.
4.  Can be combined with other clustering methods (like K-Means).


## ✗ Disadvantages

1. ! Works best with **spherical clusters** (like K-Means).
2. ✗ Doesn't always perform well with **non-uniform cluster sizes**.
3. 📖 Needs careful choice of **threshold value** for CF Tree.
4. ☹ Can lose some information during data compression.


```
from sklearn.cluster import Birch
```



## Summary Table

Feature	Description
<b>Full Form</b>	Balanced Iterative Reducing and Clustering using Hierarchies
<b>Type</b>	Hierarchical + Summarization
<b>Handles Large Data?</b>	 Yes
<b>Best for</b>	Large datasets with spherical clusters
<b>Output</b>	Cluster labels
<b>Key Parameter</b>	<code>threshold</code> (controls cluster size)

## BIRCH CLUSTERING



# 4

## Bisecting K-Means – Definition

**Bisecting K-Means** is a hierarchical version of K-Means.

It starts with **all data points in one cluster**, and then **repeatedly splits (bisects)** one cluster into two using K-Means, until the desired number of clusters is formed.





 Think of it as "**K-Means + Divide & Conquer**" method.



## ✓ Advantages



1. ✓ **Better accuracy** than regular K-Means.
2. ⚡ **Faster** than full hierarchical clustering.
3. 🧩 Works well with **large datasets**.
4. 🚀 Helps **avoid poor K-Means initialization**.
5. 📊 Creates **hierarchical cluster structure**.

## ✗ Disadvantages

1.  Still needs to specify **number of clusters (k)**.
2.  **Sensitive to outliers** (like K-Means).
3.  Can be **computationally expensive** for very large data.
4.  Only works well for **spherical-shaped clusters**.

```
from sklearn.cluster import  
BisectingKMeans
```


### BISECTING K-MEANS



Final clusters



## Summary Table

Feature	Description
Type	Hierarchical + Partitional
Base Algorithm	K-Means
Input Required	Number of clusters (k)
Best for	Large datasets with roughly spherical clusters
Output	Cluster labels + hierarchy
Speed	 Faster than standard hierarchical, slower than plain K-Means

# 5

## DBSCAN – Definition

DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a **density-based clustering** algorithm.

It groups together **closely packed points (dense regions)** and marks points that lie alone in low-density areas as **outliers**.

### **Advantages:**

1. **No need to specify number of clusters** beforehand (unlike K-Means).
2. Can find **arbitrarily shaped clusters** (not just circular).
3. Can **detect outliers/noise** easily.
4. Works well when clusters have different shapes and sizes.

## Disadvantages:

1. **Choosing  $\epsilon$  and MinPts** can be tricky.
2. Struggles with **clusters of varying density**.
3. Not suitable for **very high-dimensional data**.



```
from sklearn.cluster import DBSCAN
```

## DBSCAN

Density-Based Spatial Clustering of Applications with Noise



# 6

## Definition:

**HDBSCAN** is an **advanced version of DBSCAN** that uses **hierarchical density-based clustering**.

It automatically finds clusters of **different densities** and identifies **noise points** (outliers).

In simple terms:

- DBSCAN groups dense points but struggles when densities vary.
- **HDBSCAN** fixes that by building a **hierarchy (tree)** of clusters and selecting the most stable ones.



## **Advantages:**

- Handles **varying densities** easily
- **No need to pick eps** (automatic)
- Detects **noise/outliers**
- Finds **natural clusters**

## ✗ Disadvantages:

- Slower and **more complex** than DBSCAN
- Requires **extra library (hdbscan)**
- Harder to tune parameters

```
from sklearn.cluster import HDBSCAN
```

## HDBSCAN Clustering



Cluster 1  
(dense)

●  
Noise



Cluster 2  
(sparse)

# 7

## Definition:

**K-Means Clustering** is an **unsupervised learning algorithm** that groups data into **K number of clusters** based on how close the data points are to each other.

Each cluster has a **center point (centroid)**, and every data point belongs to the cluster with the **nearest centroid**.

## ✓ Advantages:

1. ✓ Simple and fast to use.
2. ✓ Works well with **large datasets**.
3. ✓ Easy to **understand and interpret**.
4. ✓ Works well when clusters are **clearly separated**.




## ✗ Disadvantages:


1. ✗ You must **choose K (number of clusters)** beforehand.
2. ✗ Doesn't work well with **uneven cluster sizes** or **non-spherical shapes**.
3. ✗ Sensitive to **outliers** (they can pull centroids away).
4. ✗ Different starting points can give **different results**.

```
from sklearn.cluster import KMeans
```


## K-Means Clustering



Cluster 1



Cluster 2



Cluster 3

● Centroid



# 8

## Definition:

Mean Shift is an **unsupervised clustering algorithm** that **finds clusters by locating the densest regions of data points**. It does not require you to predefine the number of clusters. The algorithm works by **shifting data points toward areas of higher density iteratively** until convergence.

## Advantages:

1. **No need to predefine clusters:** Unlike K-Means, you don't have to decide the number of clusters beforehand.
2. **Can find arbitrarily shaped clusters:** Works well with clusters that are not circular.
3. **Robust to outliers:** Less sensitive to outliers compared to some other clustering algorithms.

## Disadvantages:


1. **Computationally expensive:** Especially with large datasets, because it calculates distances repeatedly.
2. **Bandwidth selection is tricky:** The window size (bandwidth) affects results a lot; too small → too many clusters, too large → clusters merge.
3. **Not suitable for very high-dimensional data:** Performance drops when dimensions increase.

# from sklearn.cluster import Mean Shift Clustering

**Colored points** represent the data points.

**Different colors** show different clusters discovered by the algorithm.

**Red Xs** are the **cluster centers** found by Mean Shift.





# 9

## Definition

**OPTICS** stands for **Ordering Points To Identify the Clustering Structure**.

It is a **density-based clustering algorithm**, similar to DBSCAN, but more flexible.

- Unlike DBSCAN, which needs a fixed density threshold ( `eps` ) to form clusters, OPTICS can detect clusters **with varying densities**.
- Instead of producing a strict cluster assignment, it produces an **ordering of points** and a **reachability distance**, which can then be used to extract clusters at different density levels.



## Advantages

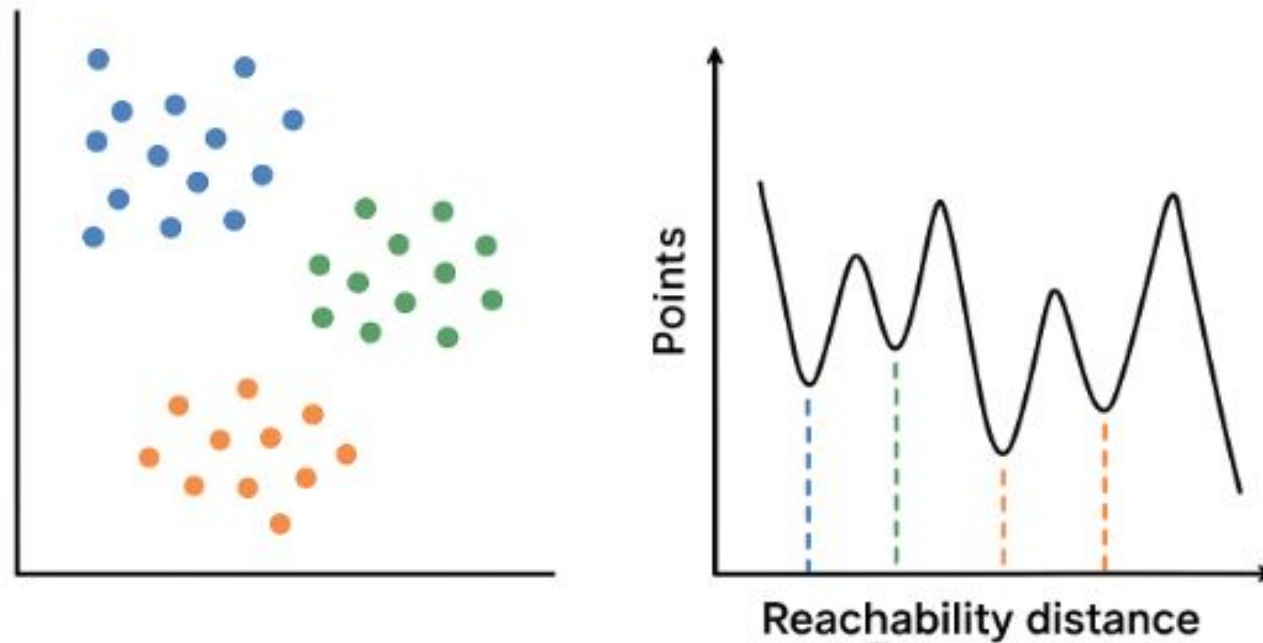
1. **Handles varying densities** – unlike DBSCAN, it doesn't require a single `eps`.
2. **Detects arbitrary shapes** – works for circular, elongated, or irregular clusters.
3. **Noise handling** – identifies outliers naturally.
4. **Flexible cluster extraction** – you can extract clusters at multiple density levels.

## Disadvantages

1. **More complex than DBSCAN** – harder to understand and implement.
2. **Slower** – especially for large datasets (complexity:  $O(n \log n)$  with indexing,  $O(n^2)$  without).
3. **Visualization required** – to extract clusters, you often need a reachability plot.
4. **Parameters still needed** – `minPts` (minimum points for a dense region) must be set correctly.

```
from sklearn.cluster import OPTICS
```

## OPTICS Clustering



# 10

## Definition

**OPTICS** stands for **Ordering Points To Identify the Clustering Structure**.

It is a **density-based clustering algorithm**, similar to DBSCAN, but more flexible.

- Unlike DBSCAN, which needs a fixed density threshold ( `eps` ) to form clusters, OPTICS can detect clusters **with varying densities**.
- Instead of producing a strict cluster assignment, it produces an **ordering of points** and a **reachability distance**, which can then be used to extract clusters at different density levels.

## Advantages:


1. Can detect **non-convex clusters** that K-Means cannot.
2. Works well with **complex cluster structures**.
3. Doesn't assume clusters are spherical.
4. Flexible: you can choose different similarity measures.

## Disadvantages:

1. Computationally expensive for **large datasets** (needs eigen decomposition).
2. Sensitive to **choice of similarity function** and parameters.
3. Number of clusters must often be specified beforehand.
4. Not very scalable to **millions of points**.

```
from sklearn.cluster import  
SpectralClustering
```

## SPECTRAL CLUSTERING



Grouping similar data points



## Definition:

The **Silhouette Score** is a measure of how well data points fit within their assigned clusters.

It tells us how **separated** and **well-formed** the clusters are.

## Typical Good Values:

- **0.7 – 1.0** → Excellent clustering
- **0.5 – 0.7** → Reasonable
- **0.25 – 0.5** → Weak
- **< 0.25** → Poor clustering



## Definition:

The **Davies–Bouldin Index** is a **measure used to evaluate clustering quality**. It checks **how well-separated** and **compact** the clusters are.

Davies–Bouldin Index	Meaning
0	Perfect clustering (clusters are very distinct)
Lower value	Better clustering
Higher value	Poor clustering (clusters overlap)

## Definition:

The **Calinski–Harabasz Index**, also called the **Variance Ratio Criterion**, is a **measure of how well the data is clustered**.

It checks how **compact** each cluster is and how **separated** the clusters are from each other.

## Range of Values:

Calinski–Harabasz Score	Meaning
Higher value	Better clustering (well-separated and compact)
Lower value	Poor clustering (overlapping or scattered clusters)

There's **no fixed range**, but **higher is always better**.

Clustering Name	Silhouette Score	Davies-Bouldin Index	Calinski-Harabasz Index	Better_Clustering?
AffinityPropagation_clustering	0.432	0.753	264.4707125	
Agglomerative_Clustering	0.553	0.578	243.0714289	
Birch_clustering	0.323	0.288	1050.99371	
BisectingKMeans_Clustering	0.446	0.823	136.2364581	
DBSCAN_Clustering	0.096	6.71	2.020445139	
HDBSCAN_Clustering	0.405	1.794	57.35299571	
MeanShift_Clustering	0.271	0.216	1378.828452	Better_Clustering
Optics_Clustering	0.292	1.623	14.29378177	
SpectralClustering	0.340	1.333	18.96466618	